


Database Management Systems


Chapter 1



What Is a Database/DBMS?

- ❖ A very large, integrated collection of data.
- ❖ Models real-world scenarios
 - Entities (e.g., students, courses)
 - Relationships (e.g., John is taking CS 5120)
- ❖ A Database Management System (DBMS) is a software package designed to store and manage databases.


Files vs. DBMS



Scenario: 100 TB of data of employees, products, salaries,


- ❖ Application must stage large datasets between main memory and secondary storage (e.g., buffering, page-oriented access, etc.)
- ❖ Special code for different queries
- ❖ Must protect data from inconsistency due to multiple concurrent users
- ❖ Crash recovery
- ❖ Security and access control

Why Use a DBMS?



- ❖ Data independence and efficient access.
- ❖ Data integrity and security.
- ❖ Uniform data administration.
- ❖ Concurrent access
- ❖ Recovery from crashes.
- ❖ Reduced application development time.


Why Study Databases??



- ❖ Shift from computation to information
 - Web applications, scientific applications, etc.
- ❖ Datasets increasing in diversity and volume.
 - Digital libraries, interactive video, Human Genome project, Earth Observation System project
 - ... need for DBMS exploding
- ❖ DBMS encompasses most of CS
 - OS, programming languages, theory, AI, multimedia, etc.


<https://www.ibm.com/analytics>

Data Models




- ❖ A data model is a collection of concepts for describing data.
- ❖ A schema is a description of a particular collection of data, using the given data model.
- ❖ The relational model of data is the most widely used model today.
 - Main concept: relation, basically a table with rows and columns.
 - Every relation has a schema, which describes the columns, or fields.

Levels of Abstraction




- ❖ Many views, single conceptual (logical) schema and physical schema.

- Views describe how users see the data.
- Conceptual schema defines logical structure
- Physical schema describes the files and indices used.




➡ Schemas are defined using DDL; data is modified/queried using DML.

Example: University Database




- ❖ Conceptual schema (Chapters 2 and 19):
 - *Students*(sid: string, name: string, login: string, age: integer, gpa: real)
 - *Courses*(cid: string, cname: string, credits: integer, fname: string)
 - *Enrolled*(sid: string, cid:string, grade: string)
 - *Faculty*(fid: string, fname: string, sal: real)
- ❖ Physical schema (Chapter 20):
 - Relations stored as unordered files.
 - Index on first column of Students.
- ❖ External Schema (View) (Chapters 3 and 25):
 - *Course_info*(cid:string, enrollment:integer, fname:string)

Data Independence




- ❖ Applications are insulated from how data is structured and stored.
- ❖ Logical data independence: Protection from changes in *logical* structure of data.
 - *Faculty*(fid: string, fname: string, sal: real)
 - *Faculty_Public*(fid: string, fname: string, office: integer)
 - *Faculty_Private*(fid: string, sal: real)
- ❖ Physical data independence: Protection from changes in *physical* structure of data.
 - Indices, file structure, disks ...

Relational Algebra



- ❖ Basic operations:
 - Selection (σ) Selects a subset of rows from a relation.
 - Projection (π) Deletes unwanted columns from relation.
 - Cross-product (\times) Allows us to combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 or in reln. 2 (or both).
- ❖ Additional operations:
 - Intersection, join, division, renaming: Not essential, but (very!) useful.
- ❖ Since each operation returns a relation, **operations can be composed!**

The SQL Query Language



- ❖ To find all 18 years old students, we can write:


```
SELECT *  
FROM Students S  
WHERE S.age=18
```

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Creating Relations in SQL



- ❖ Creates the Students relation. Observe that the type (**domain**) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.


```
CREATE TABLE Students  
(sid: CHAR(20),  
name: CHAR(20),  
login: CHAR(10),  
age: INTEGER,  
gpa: REAL)
```

- ❖ As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Enrolled  
(sid: CHAR(20),  
cid: CHAR(20),  
grade: CHAR(2))
```


<https://www.w3resource.com/sql/sql-table.php>

Concurrency Control



- ❖ Concurrent execution of user programs is essential for good DBMS performance.
 - Because disk accesses are frequent, and relatively slow, it is important to keep the cpu humming by working on several user programs concurrently.
- ❖ Interleaving actions of different user programs can lead to inconsistency: e.g., check is cleared while account balance is being computed.
- ❖ DBMS ensures such problems don't arise: users can pretend they are using a *single-user system*.

Databases make these folks happy ...




- ❖ End users and DBMS vendors
- ❖ DB application programmers
 - E.g., smart webmasters
- ❖ Database administrator (DBA)
 - Designs logical / physical schemas
 - Handles security and authorization
 - Data availability, crash recovery
 - Database tuning as needs evolve



Must understand how a DBMS works!

Structure of a DBMS

- ❖ A typical DBMS has a layered architecture.
- ❖ The figure does not show the concurrency control and recovery components.
- ❖ This is one of several possible architectures; each system has its own variations.



Summary

- ❖ DBMS used to maintain, query large datasets.
- ❖ Benefits include recovery from system crashes, concurrent access, quick application development, data integrity and security.
- ❖ Levels of abstraction give data independence.
- ❖ A DBMS typically has a layered architecture.
- ❖ DBAs hold responsible jobs and are **well-paid!** 😊
- ❖ DBMS R&D is one of the broadest, most exciting areas in CS.

