


Introduction to Database Design

Chapter 2



Overview of Database Design


❖ Requirements Analysis

- The very first step in designing a database application.
- We focus on the design of the database.

❖ Conceptual design: (*ER Model is used at this stage.*)


- What are the *entities* and *relationships* in the enterprise?
- What information about these entities and relationships should we store in the database?
- What are the *integrity constraints* or *business rules* that hold?

Overview of Database Design (Cont.)




sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@ee	18	3.2

Overview of Database Design (Cont.)



- A database 'schema' in the ER Model can be represented pictorially (*ER diagrams*).
- ❖ Logical Database Design
 - Choose a DBMS to implement our database design.
 - Map an **ER diagram** into a **relational schema** (Chapter 3).

Overview of Database Design (Cont.)



❖ Schema Refinement (Chapter 19)

- Analyze the current schema to identify potential problems and to refine it (performance criteria).
- Theory – *normalizing relations*.


❖ Physical Database Design (Chapter 20)

- Build indices on some tables and clustering some tables.

❖ Application and Security Design (Chapter 21)

- Consider aspects of the application that go beyond the database itself.

ER Model Basics




❖ Entity: Real-world object distinguishable from other objects. An entity is described (in DB) using a set of attributes.

❖ Entity Set: A collection of similar entities.
E.g., all employees.

- All entities in an entity set have the same set of attributes. (Until we consider ISA hierarchies, anyway!)
- Each entity set has a *key* (primary key - SSN).
- Each attribute has a *domain*.

ER Model Basics (Contd.)



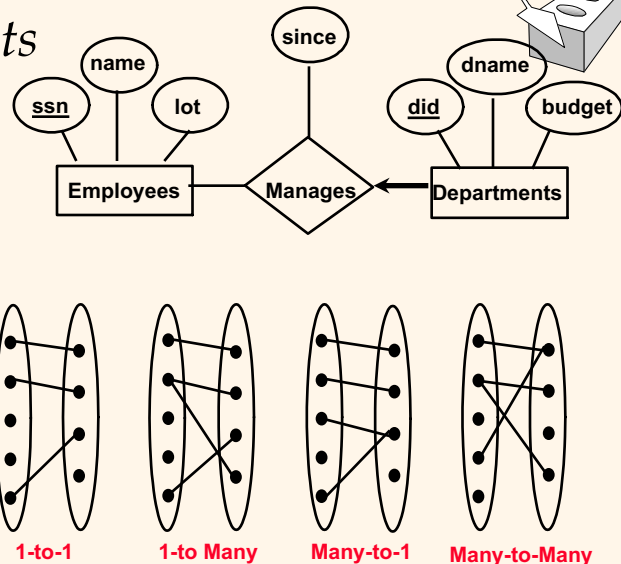
- ❖ Relationship: Association among two or more entities.
E.g., Barbara works in the CS department.
- ❖ Relationship Set: Collection of similar relationships.
 - An n-ary relationship set R relates n entity sets E1 ... En; each relationship in R involves entities $e1 \in E1, \dots, en \in En$
 - Same entity set could participate in different relationship sets, or in different “roles” in same set.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

7

Key Constraints

- ❖ Consider Works_In:
An employee can work in many departments; a dept can have many employees.
- ❖ In contrast, each dept has **at most** one manager, according to the key constraint on Manages.




Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

8

Participation Constraints

❖ Does every department have a manager?


- If so, this is a *participation constraint*: the participation of Departments in Manages is said to be *total* (vs. *partial*).
 - Every Departments entity must appear in an instance of the Manages relationship.



Weak Entities

❖ A *weak entity* can be identified uniquely only by considering the primary key of another (*owner*) entity.

- Owner entity set and weak entity set must participate in a *one-to-many* relationship set (one owner, many weak entities).
- Weak entity set must have *total participation* in this *identifying relationship set*.
- Identifying owner (the way to identify a weak entity) – SSN + pname (*partial key*).



ISA ('is a') Hierarchies

❖ As in C++, or other PLs, attributes are inherited.

❖ If we declare A **ISA** B, every A entity is also considered to be a B entity.

❖ Reasons for using ISA:

- To add descriptive attributes specific to a subclass.
- To identify entities that participate in a relationship.

❖ **Overlap constraints**: Can Joe be an Hourly_Emps as well as a Contract_Emps entity? (*Allowed/disallowed*) *No by default*

❖ **Covering constraints**: Does every Employees entity also have to be an Hourly_Emps or a Contract_Emps entity? (*Yes/no*) *No by default*

