



Overview of Query Evaluation

Chapter 12



Overview of Query Evaluation


- ❖ Plan: Tree of R.A. ops, with choice of alg for each op.
 - Each operator is typically implemented using a 'pull' interface: when an operator is 'pulled' for the next output tuples, it 'pulls' on its inputs and computes them.
- ❖ Two main issues in query optimization:
 - For a given query, **what plans are considered?**
 - Algorithm to search **plan space** for the cheapest (estimated) plan.
 - How is the **cost of a plan estimation?**
- ❖ **Ideally**: Want to find the best plan. **Practically**: Avoid worst plans!

Overview of Query Evaluation



```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```

RA Tree:




Cost Estimation



- ❖ For each plan considered, must estimate cost:
 - Must **estimate cost** of each operation in plan tree.
 - Depends on input cardinalities.
 - Must also **estimate size of result** for each operation in tree!
 - Use information about the input relations.

Relational Algebra



- ❖ Basic operations:
 - Selection (σ) Selects a subset of rows from a relation.
 - Projection (π) Deletes unwanted columns from relation.
 - Cross-product (\times) Allows us to combine two relations.
 - Set-difference ($-$) Tuples in reln. 1, but not in reln. 2.
 - Union (\cup) Tuples in reln. 1 or in reln. 2 (or both).
- ❖ Additional operations:
 - Intersection, join, division, renaming: Not essential, but (very!) useful.
- ❖ Since each operation returns a relation, **operations can be composed!**

Statistics and Catalogs



- ❖ Need information about the relations and indexes involved. **Catalogs** typically contain at least:
 - # tuples (NTuples) and # pages (NPages) for each relation.
 - # distinct key values (NKeys) and NPages for each index.
 - Index height, low/high key values (Low/High) for each tree index.
- ❖ Catalogs updated periodically.
 - Updating whenever data changes is too expensive; lots of approximation anyway, so slight inconsistency ok.

Size Estimation and Reduction Factors

```
SELECT attribute list
FROM relation list
WHERE term1 AND ... AND termk
```

- ❖ Consider a query block:
- ❖ Maximum # tuples in result is the product of the cardinalities of relations in the FROM clause.
- ❖ *Reduction factor (RF)* associated with each *term* reflects the impact of the *term* in reducing result size. *Result cardinality* = Max # tuples * product of all RF's.
 - Implicit assumption that *terms* are independent!
 - Term *col=value* has RF $1/NKeys(I)$, given index *I* on *col*
 - Term *col1=col2* has RF $1/MAX(NKeys(I1), NKeys(I2))$
 - Term *col>value* has RF $(High(I)-value)/(High(I)-Low(I))$

Some Common Techniques

- ❖ Algorithms for evaluating relational operators use some simple ideas extensively:
 - **Indexing:** Can use WHERE conditions to retrieve small set of tuples (selections, joins)
 - **Iteration:** Sometimes, faster to scan all tuples even if there is an index. (And sometimes, we can scan the data entries in an index instead of the table itself.)
 - **Partitioning:** By using sorting or hashing, we can partition the input tuples.

** Watch for these techniques as we discuss query evaluation!*

Access Path



- ❖ An access path is a method of retrieving tuples:
 - File scan, or index that matches a selection (in the query)
- ❖ A tree index matches (a conjunction of) terms that involve only attributes in a *prefix* of the search key.
 - E.g., Tree index on $\langle a, b, c \rangle$ matches the selection $a=5$ AND $b=3$, and $a=5$ AND $b>6$, but not $b=3$.
- ❖ A hash index matches (a conjunction of) terms that has a term *attribute = value* for every attribute in the search key of the index.
 - E.g., Hash index on $\langle a, b, c \rangle$ matches $a=5$ AND $b=3$ AND $c=5$; but it does not match $b=3$, or $a=5$ AND $b=3$, or $a>5$ AND $b=3$ AND $c=5$.

One Approach to Selections



- ❖ Find the *most selective access path*, retrieve tuples using it, and apply any remaining terms that do not match the index:
 - *Most selective access path*: An index or file scan that we estimate will require the fewest page I/Os.
 - Terms that match this index reduce the number of tuples retrieved; other terms are used to discard some retrieved tuples, but do not affect number of tuples/pages fetched.
 - Consider $day<8/1/2016$ AND $bid=5$ AND $sid=3$. A B+ tree index on *day* can be used; then, $bid=5$ and $sid=3$ must be checked for each retrieved tuple. Similarly, a hash index on $\langle bid, sid \rangle$ could be used; $day<8/1/2016$ must then be checked.

Relational Operation - Selection



- ❖ No indexes => scan the file
- ❖ With an index, cost depends on #qualifying tuples, and **clustering**.
 - Cost of finding qualifying data entries (typically small) plus cost of retrieving records (could be large w/o clustering).
 - In example, assuming uniform distribution of names, about 10% of tuples qualify (100 pages, 10000 tuples). With a clustered B+ tree index, cost is little more than 100 I/Os; if unclustered, up to 10000 I/Os!

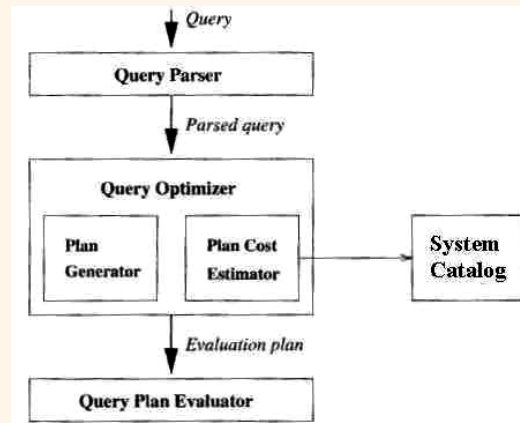

Projection

```
SELECT  DISTINCT
        R.sid, R.bid
FROM    Reserves R
```



- ❖ The expensive part is removing duplicates.
 - SQL systems don't remove duplicates unless the keyword DISTINCT is specified in a query.
- ❖ *Sorting Approach*: Sort on <sid, bid> and remove duplicates. (Can optimize this by dropping unwanted information while sorting.)
- ❖ *Hashing Approach*: Hash on <sid, bid> to create partitions. Load partitions into memory one at a time and eliminate duplicates.
- ❖ If there is an index with both R.sid and R.bid in the search key, may be cheaper to *sort data entries in the index*.

Query Optimizer



Schema for Examples




Sailors (sid: integer, sname: string, rating: integer, age: real)
Reserves (sid: integer, bid: integer, day: dates, rname: string)

- ❖ Similar to old schema; *rname* added for variations.
- ❖ Reserves:
 - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- ❖ Sailors:
 - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

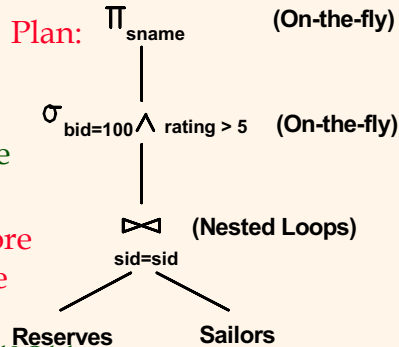
Motivating Example

```
SELECT S.sname
FROM Reserves R, Sailors S
WHERE R.sid=S.sid AND
      R.bid=100 AND S.rating>5
```


RA Tree:



- ❖ Cost: $1000 + 1000 * 500$ I/Os
- ❖ By no means the worst plan!
- ❖ Misses several opportunities: selections could have been 'pushed' earlier, no use is made of any available indexes, etc.
- ❖ **Goal of optimization:** To find more efficient plans that compute the same answer.



Alternative Plan



- ❖ **Main difference:** push selects.
- ❖ With 5 buffers, **cost of plan:**
 - Scan Reserves (1000) + write temp T1 (10 pages, if we have one hundred boats, uniform distribution).
 - Scan Sailors (500) + write temp T2 (250 pages, if we have 10 ratings).
 - Sort T1 ($2 * 2 * 10$), sort T2 ($2 * 4 * 250$), merge ($10 + 250$)
 - **Total: 4060 page I/Os.**
- ❖ **If we 'push' projections,** T1 has only *sid*, T2 only *sid* and *sname*:
 - The cost of the join step drops significantly, **total < 2000.**

Summary



- ❖ There are several alternative evaluation algorithms for each relational operator.
- ❖ A query is evaluated by converting it to a tree of operators and evaluating the operators in the tree.
- ❖ Must understand query optimization in order to fully understand the performance impact of a given database design (relations, indexes) on a workload (set of queries).
- ❖ Two parts to optimizing a query:
 - Consider a set of alternative plans.
 - Must prune search space; typically, left-deep plans only.
 - Must estimate cost of each plan that is considered.
 - Must estimate size of result and cost for each plan node.
 - *Key issues*: Statistics, indexes, operator implementations.